

Aayush Aashish Mehta

+1 469-901-3435 | aayushmehta0924@gmail.com | [LinkedIn](#) | [Portfolio](#)

PROFESSIONAL SUMMARY

- 3+ years of experience in BI engineering, data engineering, and AI-powered analytics, specializing in last-mile pipeline development, Power BI reporting, and Azure-native data solutions for enterprise-scale systems.
- Expert in building end-to-end ETL/ELT pipelines using Python, PySpark (Spark), Azure Data Factory (ADF), and Databricks, delivering data to Power BI, Power Automate, and Power Apps across diverse enterprise sources.
- Proficient in Power BI development with advanced DAX (time intelligence, context transitions, variables), star-schema modeling, Row-Level Security (RLS), and Power BI Service publishing for cross-functional stakeholders.
- Experienced integrating AI and LLM capabilities (Azure OpenAI, Azure AI Studio, Semantic Kernel) into data pipelines and business applications via APIs, batch jobs, and event-driven Azure Functions workflows.
- Hands-on with Databricks (SDK, Jobs API) and Medallion Architecture (Bronze/Silver/Gold) for scalable lakehouse design using Azure Synapse Analytics and Microsoft Fabric.
- Built Power Automate cloud flows and Power Apps canvas applications to automate business processes, report distribution, and data-driven workflows, improving operational efficiency across teams.
- Implemented robust data quality and governance frameworks covering schema validation, row counts, null thresholds, and data-type consistency across SAP HANA, Salesforce, and ServiceNow pipelines.
- Experienced integrating enterprise ERP data (SAP HANA, Salesforce, ServiceNow, Workday) into unified Azure Synapse analytical models, Power BI dashboards, and AI-powered applications.
- Built and deployed Azure ML pipelines for predictive analytics and churn prediction, integrating risk scores into Power BI KPI dashboards and CRM workflows for AI-powered business actions.
- Proficient in CI/CD using Azure DevOps (YAML pipelines, Databricks SDK deployments) for automated pipeline and Power BI report promotion across DEV, QA, and PROD environments.
- Designed multi-agent AI systems (Azure AI Studio, Semantic Kernel, Azure AI Agent Service) enabling natural-language querying of enterprise data and automated DAX/SQL generation into Power BI workspaces.
- Proven ability to train and support stakeholders on BI tools; delivered Power BI workshops for 15+ staff, driving self-service analytics adoption and reducing analyst support tickets by ~30%.
- Adept at translating complex business requirements into modular, reusable ADF pipeline templates and Databricks SDK scripts; documented pipelines, DAX logic, and dashboard guides for cross-team onboarding.

TECHNICAL SKILLS

Programming Languages: Python, SQL, DAX (Advanced), PySpark (Spark), Shell Scripting, JavaScript, C/C++

BI & Visualization: Power BI (Dashboards, DAX, Star-Schema Modeling, RLS, Power BI Service), Power Automate, Power Apps, Looker, Tableau, Matplotlib, Seaborn

Data Engineering & Processing: Azure Data Factory (ADF), Databricks (SDK, Jobs API, Notebooks), Apache Airflow, dbt, Apache Kafka, Docker, ETL/ELT, Medallion Architecture (Bronze/Silver/Gold)

Azure Services: Azure Synapse Analytics, Microsoft Fabric Lakehouse, Azure AI Studio, Azure OpenAI (GPT-4o), Azure ML, Azure AI Agent Service, Azure AI Search, Semantic Kernel, Azure DevOps, Azure Event Hubs, Azure Monitor, Azure Functions, Logic Apps, ADLS Gen2

AI & LLM Integration: Azure OpenAI, Semantic Kernel, Azure AI Agent Service, LangChain, LlamaIndex, Vertex AI, Scikit-learn, TensorFlow, Keras, PyTorch

Cloud Platforms: Microsoft Azure, GCP (BigQuery, Vertex AI, Cloud Composer, Pub/Sub, Cloud Functions, Dataproc), AWS (S3, Redshift)

ERP & Enterprise Systems: SAP HANA, SAP BusinessObjects, Salesforce, ServiceNow, Workday

Databases: Azure Synapse, Microsoft Fabric, BigQuery, Amazon Redshift, PostgreSQL, MySQL, SQL Server, MongoDB, Neo4j, Oracle, SAP HANA

Data Modeling: Star/Snowflake Schema, Medallion Architecture, Dimensional Modeling, ERwin

Version Control & CI/CD: Azure DevOps, Git, GitLab, Databricks SDK (CLI + Jobs API)

Web Development: Flask, Node.js, React.js, MERN Stack

Methodologies: Agile (Scrum), SDLC

PROFESSIONAL EXPERIENCE

ITDataBuild

BI Engineer

Flower Mound, TX | Jan 2026 – Present

- Engineered end-to-end ETL/ELT pipelines using Python, SQL, PySpark (Spark), and Azure Data Factory (ADF) to extract and transform data from SAP HANA source systems, loading into Azure Synapse Analytics and Microsoft Fabric Lakehouse (Bronze/Silver/Gold Medallion layers) for Power BI consumption.
- Built Power BI semantic models (star-schema) on Synapse and Fabric datasets, authoring advanced DAX, time intelligence, rolling 13-week averages, YoY variance, dynamic ranking; powering executive KPI dashboards for finance, operations, and supply chain teams.
- Developed 12+ Power BI reports and dashboards with Row-Level Security (RLS), scheduled refresh via Power BI Service, and drill-through navigation, reducing stakeholder ad-hoc reporting requests by ~45%.
- Architected a multi-agent AI pipeline using Azure AI Studio and Azure OpenAI with Semantic Kernel orchestration that interprets natural-language business questions, auto-generates validated DAX and T-SQL queries, and publishes results directly into Power BI workspaces; cutting analyst turnaround from days to under an hour.
- Built Power Automate cloud flows and a Power Apps canvas application enabling operations teams to trigger ADF pipelines, submit data requests, and view embedded Power BI reports — eliminating manual follow-ups across 4 business teams and improving self-service data access.
- Implemented data quality and governance checks within ADF pipelines on row counts, schema parity, null thresholds, data-type consistency with automated alerting via Azure Monitor and Logic Apps.
- Set up Azure DevOps CI/CD with YAML templates for automated deployment of ADF pipeline definitions, Power BI datasets (PBIP format), and Databricks jobs and workflows using the Databricks SDK (CLI + Jobs API) across DEV, QA, and PROD, enforcing zero-downtime releases.

- Automated 6 legacy Python batch jobs into parameterized ADF pipelines with retry policies, dependency chains, and failure alerting via Azure Monitor improving pipeline reliability by ~60%.
- Integrated Azure Event Hubs for near-real-time ingestion of operational telemetry into Fabric Lakehouse, enabling streaming dashboard updates in Power BI and replacing batch-only reporting for time-sensitive operations.
- Developed reusable ADF pipeline templates and Databricks SDK-based deployment scripts that cut new data source onboarding time by ~50%, accelerating delivery across subsequent migration sprints.

Sabre

Flower Mound, TX | May 2025 – Dec 2025

Data / AI Engineer

- Built last-mile ETL pipelines using Azure Data Factory, Databricks, and PySpark (Spark) to integrate SAP, ServiceNow, and Salesforce data into a Customer 360 model in Azure Synapse Analytics, consolidated for Power BI consumption across finance, sales, and operations.
- Developed AES-Assist, a multi-agent AI system using Azure AI Studio, Semantic Kernel, and Azure AI Agent Service that interprets stakeholder questions, auto-generates Power BI DAX and Synapse SQL, and surfaces answers via a Microsoft Teams bot — reducing ad-hoc analytics requests by ~40%.
- Integrated an Azure ML churn prediction model into the Customer 360 Databricks pipeline using PySpark; scoring accounts at batch ingestion time, writing risk tiers back to Synapse, and surfacing predicted churn probabilities as a live KPI in Power BI dashboards for the sales team.
- Integrated Azure OpenAI APIs and event-driven Azure Functions into business workflows, enabling AI-powered automation for report summarization, anomaly flagging, and proactive insight delivery.
- Computed and standardized 15+ KPIs (Account Engagement Score, churn risk, pipeline velocity, NPS) in Synapse, making them queryable across Power BI, finance, sales, and operations teams.
- Designed and built 8 Power BI dashboards with advanced DAX, drill-through pages, and RLS including a 6-KPI Account Engagement Score dashboard tracking historical trends for global enterprise clients.
- Created a Text-to-SQL assistant over Synapse Customer 360 tables using Azure OpenAI and Azure AI Search, enabling stakeholders to query KPIs in plain English with no SQL knowledge.
- Built and maintained parameterized ADF pipelines for incremental batch ingestion of Salesforce, ServiceNow, and SAP data into Synapse, with automated schema drift detection and Azure Monitor alerting.
- Migrated SQL transformation logic to dbt on Databricks with Medallion Architecture (staged/intermediate/mart layers), Slim CI on pull requests, and model documentation tied to Power BI lineage views.
- Delivered 2 Power BI training workshops for 15 staff members covering end-to-end dashboard development, DAX fundamentals, Power BI Service publishing thereby boosting self-service adoption and cutting analyst support tickets by ~30%.
- Built Power Automate cloud flows to automate Power BI report subscriptions, alert-triggered notifications, and data approval workflows, replacing manual stakeholder coordination across 3 business teams.

Analytics IT

Sept 2022 – June 2023

Data / AI Engineer

- Designed and deployed automated ETL pipelines using Apache Airflow and BigQuery to support customer segmentation for a global retail client, enabling last-mile delivery to Power BI dashboards.
- Built a data lakehouse using Medallion Architecture by integrating HR, Salesforce, and ServiceNow data; consolidated 18+ business KPIs and developed interactive Power BI dashboards for cross-functional reporting.
- Standardized transformations in DBT: built staged/intermediate/mart layers for HR, Salesforce, and ServiceNow; incremental models partitioned by event_date; and added data quality tests (unique/not_null/relationships).
- Built an NLP chatbot for a lead enrichment platform, enabling marketing teams to access campaign metrics and client profiles via conversational voice or text queries.
- Created an ML pipeline using Vertex AI to predict churn probabilities, integrating the output with Salesforce CRM for personalized AI-powered retention workflows.
- Developed Power BI and Looker dashboards to visualize campaign ROI, customer funnel health, and cohort drop-off patterns for business stakeholders.
- Engineered modular DAGs for batch ingestion of Google Ads campaign metadata into BigQuery for attribution modeling and reporting.
- Tuned SQL transformation logic in BigQuery to improve dashboard load times by 35% and reduce query costs.
- Designed schema monitoring scripts to track upstream changes and auto-notify analysts on pipeline breakage risks, enforcing data governance standards.

Nividata Consultancy

Ahmedabad, India | May 2021 – Aug 2021

Data Engineer

- Led development of batch ETL pipelines using Python and Apache Airflow to process product and sales data for an e-commerce client, delivering output to Power BI and Tableau dashboards.
- Integrated AWS S3 as a centralized staging layer and built data ingestion workflows for MySQL and MongoDB sources.
- Developed a data validation layer using Python and Pandera to enforce schema integrity and data governance standards before loading into Redshift.
- Built interactive Power BI and Tableau dashboards tracking sales performance, conversion rates, and product return ratios across regions.
- Implemented Python scripts to detect and alert anomalies in weekly sales trends, improving revenue monitoring responsiveness.
- Migrated legacy Excel-based reporting workflows to automated Power BI dashboards, cutting manual reporting effort by over 50%.
- Documented all ETL processes and transformation logic in Confluence, enabling team-wide transparency and maintainability.

Bosky Buildcon

May 2020 – Dec 2020

Data Analyst

- Conducted cost-efficiency analysis across 12 active construction projects using SQL and Excel, identifying a 15% overhead savings opportunity through material reallocation.
- Built interactive visual reports using Matplotlib and Seaborn to track project burn rates, workforce allocation, and vendor performance.
- Automated daily extraction of site data from internal tools using Python, reducing manual reporting workload by 40%.
- Created Tableau dashboards for site supervisors to monitor budget utilization and forecast material demand.
- Standardized data entry workflows and implemented validation scripts to reduce reporting inconsistencies across project sites.

ACTIVITIES

Youth India Foundation

Nov 2021 – Nov 2022

Head of Finance

- Managed budgeting, expense tracking, and revenue planning for state and national events, ensuring financial transparency and accountability.
- Used Excel to analyze event profitability, visualize revenue trends, and prepare weekly financial reports for leadership.
- Collaborated with event and operations teams, using data insights to optimize resource allocation and reduce overall event costs.

TwoWaits

Nov 2021 – Dec 2021

Web Developer

- Developed a secure email authentication module in Node.js and integrated it with a Flutter client application for seamless user onboarding.
- Contributed to frontend improvements in the Flutter app, enhancing UI responsiveness and refining input validation during login and signup flows.

PUBLICATION

DeSculpt: Indian Temple Sculpture Iconography

- Co-authored and published a deep learning framework for 9-class classification of Navagraha deity sculptures, addressing extreme data scarcity via data augmentation and transfer learning on ResNet50, VGG16, and MobileNet.
- Achieved 91% classification accuracy; documented dataset creation, preprocessing pipeline, architecture comparison, and evaluation metrics in the published paper.

EDUCATION

Master of Science in Computer Science

Arizona State University, USA | CGPA: 3.73/4 | Awarded NAMU Scholarship

Bachelor of Technology in Computer Science Engineering

S.R.M. University, Chennai, India | CGPA: 9.38/10